

# HPSC Lab2

Luna McBride, Mehmood Ali

September 13 2023

## 1 Parallel Design

This assignment utilizes multiple cores to build out a grid. The objective is to get us thinking about how parallel pieces should connect to take advantage of the strengths present in parallel computing. As a result, this task functions more as a "change your way of thinking" assignment rather than a difficult coding assignment like those present in other classes. Nonetheless, parallelizing a grid like this turned out a lot harder than it initially seemed.

In this lab, we need to implement a finite difference solver for Laplace's equation on a grid that is decomposed among multiple processors using MPI for parallelism. The linear system is solved iteratively, and the code handles boundary conditions and communication between processors effectively.

- **fd.cpp**: This is the main program file that contains the main function. It initializes MPI, performs grid decomposition, forms a linear system, solves it using either Gauss-Seidel or Jacobi iteration, and plots the results. It also includes a class called `LaplacianOnGrid`.
- **fd.h**: This header file contains various include statements, macros, and typedefs used throughout the code. It's included in both **fd.cpp** and **gauss\_seidel.h**.
- **gauss\_seidel.h**: This header file defines a function **GS\_or\_Jacobi** that performs Gauss-Seidel or Jacobi iterations to solve a linear system. It's used in **fd.cpp** for solving the linear system.
- **mpiinfo.h**: This header file defines the **mpiInfo** class, which contains information about the MPI environment, such as the processor's rank, number of processors, grid decomposition, and functions for exchanging boundary information between processors.

In this lab we use a parallel approach for tackling the Laplace equation through a finite difference method using MPI (Message Passing Interface). It begins by dividing the problem domain into a grid, with each grid point representing a variable to be solved for. This 2D grid is then distributed across multiple MPI processes, creating a structured processor grid, where each process is assigned a unique identifier (myPE) and a position in the processor grid (iPE, jPE). Communication plays a pivotal role, as each processor needs to exchange boundary values with its neighboring processors to facilitate the solution of the Laplace equation. These neighbors are determined by the processor's position within the grid. The communication itself is handled through MPI point-to-point communication routines. The formulation of the linear system is central to the code's functionality.

The **LaplacianOnGrid** class defines the grid and its associated variables, ultimately constructing the matrix  $A$  and the right-hand side vector  $b$  that constitute the linear system  $Ax = b$ , representing the Laplace equation. The code accommodates various boundary conditions, including Dirichlet boundary conditions, depending on the specific problem requirements. For the iterative solution phase, the code implements two iterative methods, namely **Gauss-Seidel** and **Jacobi**. These methods continue to iterate until a convergence criterion is met. Each iteration involves not only updating the interior grid points but also communicating boundary values with neighboring processors and verifying convergence against a specified tolerance or a maximum iteration count. Parallel communication, managed by the **ExchangeBoundaryInfo** function, is vital for synchronizing boundary values between neighboring processors. This function sends values from one processor to another and receives values in return. The received boundary values are then employed to update boundary nodes and apply relevant boundary conditions for the subsequent iteration.

The main parallelization strategy centers on domain decomposition. By dividing the computational domain into smaller subdomains, each assigned to a different MPI process. Processors exclusively focus on their designated portions of the grid and interact solely with neighboring processors to solve the problem.

## 2 Self Evaluation

### 2.1 Luna

I really had a problem with the loading and compiling more than anything else. I have included the ex01.bat slurm file I had to modify, but it was throwing errors with required library dependencies until I made the change to update and use gcc instead of intel like originally written. On top of that, I made an error in changing around nPE<sub>x</sub> and nPE<sub>y</sub> that I could not even see until libstdc++ was fixed and various loading attempts occurred. I really spent a lot more time fighting my computer than actually coding the mpiInfo. I am glad Mehmood was able to get that part done so quickly and efficiently in that time.

### 2.2 Mehmood

At first, I had trouble running my program on a single processor. It took me a while to figure out that I needed to make changes to the slurm file by setting nPE<sub>x</sub> and nPE<sub>y</sub> to 1. I also had to modify the task nodes. After some trial and error, I finally managed to run it successfully.

I also had some issues when trying to plot graphs with Gnuplot. I made a simple mistake where every time I ran the command "load pc," it asked for a file name. It took me a while to realize that I should use the command "load 'pc'" with single quotes around 'pc'.

Implementing boundary conditions and debugging the program was a bit challenging. Debugging took up a significant amount of time. I used multiple "cout" statements to check the program's outputs. When working on the boundary conditions, I learned that it was crucial to make sure that neighboring processors communicated accurately and on time. In the beginning, I got some strange plots due to communication errors between processors. I spent most of my time debugging for this lab. Due to the complex nature of the problem, we were solving, it took some time, but eventually, after watching the lecture multiple times, I was able to make it work.

### 3 Appendix A: Parallel Code

```

1  #include "fd.h"
2
3  void Exit()
4  {
5      MPI_Barrier(MPI_COMM_WORLD);
6      MPI_Finalize();
7      exit(0);
8  }
9
10
11 // ==
12 // ||
13 // ||   C L A S S :   m p i I n f o
14 // ||
15 // ==
16
17 class mpiInfo
18 {
19 public:
20
21     int myPE;
22     int numPE;
23     int nRealx, nRealy;
24     int nPEEx, nPEy;
25     int iPE , jPE;
26     int iMin, iMax, jMin, jMax ; // The global i-j numbers on this processor
27     int nei_n, nei_s, nei_e, nei_w;
28     int countx, county;
29
30     double *phiL, *phiR;
31     double *phiT, *phiB;
32
33     double *phiSend_n, *phiSend_s;
34     double *phiSend_e, *phiSend_w;
35     double *phiRecv_n, *phiRecv_s, *phiRecv_e, *phiRecv_w;
36
37     MPI_Status status;
38     int err;
39     int tag;
40     MPI_Request request;
41
42     // -
43     // |
44     // |   GridDecomposition: Set up PE numbering system in figure below and
45     // |                           establish communication arrays.
46     // |
47     // |                           nPEEx -- number of PEs in the x-direction
48     // |                           nPEy  -- number of PEs in the y-direction
49     // |                           numPE = total number of PEs
50     // |
51     // |
52     // |
53     // |

```

+-----+	+-----+	. . .	+-----+
			numPE

```

54 // | | | | |
55 // | +-----+-----+ . . . +-----+
56 // | . . . . .
57 // | . . . . .
58 // | . . . . .
59 // | +-----+-----+ . . . +-----+
60 // | | nPEx | nPEx+1 | | |
61 // | | | | | | |
62 // | | | | | | |
63 // | +-----+-----+ . . . +-----+
64 // | | 0 | 1 | | nPEx-1 |
65 // | | | | | | |
66 // | | | | | | |
67 // | +-----+-----+ . . . +-----+
68 // |
69 // |
70 // -
71
72
73 void GridDecomposition(int _nPEx, int _nPEy, int nCellx , int nCelly)
74 {
75
76     nRealx = nCellx + 1;
77     nRealy = nCelly + 1;
78
79     // Store and check incoming processor counts
80
81     nPEx = _nPEx;
82     nPEy = _nPEy;
83
84     if (nPEx*nPEy != numPE)
85     {
86         if ( myPE == 0 ) cout << "Fatal Error:  Number of PEs in x-y directions do
not add up to numPE" << endl;
87         MPI_Barrier(MPI_COMM_WORLD);
88         MPI_Finalize();
89         exit(0);
90     }
91
92     // Get the i-j location of this processor, given its number.  See figure above
93     :
94
95     jPE = int(myPE/nPEx);
96     iPE = myPE - jPE*nPEx;
97
98     // Set neighbor values
99
100     nei_n = nei_s = nei_e = nei_w = -1;
101
102     if ( iPE > 0 )
103     {
104         nei_w = myPE - 1;
105     }
106     if ( jPE > 0 )
107     {

```

```

107     nei_s = myPE-2;
108     }
109     if ( iPE < nPEx-1 )
110     {
111     nei_e = myPE+1;
112     }
113     if ( jPE < nPEy-1 )
114     {
115     nei_n=myPE+2;
116     }
117
118
119     countx = nRealx + 2;
120     county = nRealy + 2;
121
122     phiL = new double [ county ];
123     phiR = new double [ county ];
124     phiT = new double [ countx ];
125     phiB = new double [ countx ];
126
127     phiSend_n = new double [ countx ];
128     phiSend_s = new double [ countx ];
129     phiSend_e = new double [ county ];
130     phiSend_w = new double [ county ];
131     phiRecv_n = new double [ countx ];
132     phiRecv_s = new double [ countx ];
133     phiRecv_e = new double [ county ];
134     phiRecv_w = new double [ county ];
135
136     tag = 0;
137 }
138
139 void ExchangeBoundaryInfo(VD &Solution, VD &b)
140 {
141     sLOOP phiSend_n[s] = 0.;
142     sLOOP phiSend_s[s] = 0.;
143     tLOOP phiSend_e[t] = 0.;
144     tLOOP phiSend_w[t] = 0.;
145
146     // -----
147     // (1) Parallel communication on PE Boundaries:    ** See fd.h for tLOOP and
148     // sLOOP macros **
149     // -----
150
151     // (1.1) Put them into communication arrays
152
153     sLOOP phiSend_n[s] = Solution[pid(s, nRealy - 1)];
154     sLOOP phiSend_s[s] = Solution[pid(s, 2)];
155     tLOOP phiSend_e[t] = Solution[pid(nRealx - 1, t)];
156     tLOOP phiSend_w[t] = Solution[pid(2, t)];
157
158     // (1.2) Send them to neighboring PEs
159     if ( nei_n >= 0 ) err = MPI_Isend(phiSend_n, nRealx + 2, MPI_DOUBLE, nei_n, 0,
160     MPI_COMM_WORLD, &request);
161     if ( nei_s >= 0 ) err = MPI_Isend(phiSend_s, nRealx + 2, MPI_DOUBLE, nei_s, 0,

```

```

160 MPI_COMM_WORLD, &request);
161 if ( nei_e >= 0 ) err = MPI_Isend(phiSend_e, nRealy + 2, MPI_DOUBLE, nei_e, 0,
    MPI_COMM_WORLD, &request);
162 if ( nei_w >= 0 ) err = MPI_Isend(phiSend_w, nRealy + 2, MPI_DOUBLE, nei_w, 0,
    MPI_COMM_WORLD, &request);
163
164 // (1.3) Receive values from neigboring PEs' physical boundaries.
165
166 if ( nei_n >= 0 ) { err = MPI_Irecv(phiRecv_n, nRealx + 2, MPI_DOUBLE, nei_n,
    0, MPI_COMM_WORLD, &request); MPI_Wait(&request,&status); }
167 if ( nei_s >= 0 ) { err = MPI_Irecv(phiRecv_s, nRealx + 2, MPI_DOUBLE, nei_s,
    0, MPI_COMM_WORLD, &request); MPI_Wait(&request,&status); }
168 if ( nei_e >= 0 ) { err = MPI_Irecv(phiRecv_e, nRealy + 2, MPI_DOUBLE, nei_e,
    0, MPI_COMM_WORLD, &request); MPI_Wait(&request,&status); }
169 if ( nei_w >= 0 ) { err = MPI_Irecv(phiRecv_w, nRealy + 2, MPI_DOUBLE, nei_w,
    0, MPI_COMM_WORLD, &request); MPI_Wait(&request,&status); }
170
171 // (1.4) If new information was received, store it in the candy-coating values
172
173 if ( nei_n >= 0 ) sLOOP Solution[pid(s, nRealy + 1)] = phiRecv_n[s] ;
174 if ( nei_s >= 0 ) sLOOP Solution[pid(s, 0)] = phiRecv_s[s] ;
175 if ( nei_e >= 0 ) tLOOP Solution[pid(nRealx + 1, t)] = phiRecv_e[t] ;
176 if ( nei_w >= 0 ) tLOOP Solution[pid(0, t)] = phiRecv_w[t] ;
177
178 // (1.5) Apply exchanged information as BCs
179
180 if ( nei_n >= 0 ) sLOOP b[pid(s, nRealy + 1)] = phiRecv_n[s] ;
181 if ( nei_s >= 0 ) sLOOP b[pid(s, 0)] = phiRecv_s[s] ;
182 if ( nei_e >= 0 ) tLOOP b[pid(nRealx + 1, t)] = phiRecv_e[t] ;
183 if ( nei_w >= 0 ) tLOOP b[pid(0, t)] = phiRecv_w[t] ;
184 };
185
186 int pid(int i,int j) { return (i+1) + (j)*(nRealx+2); }
187
188 };

```

Listing 1: mpiInfo.h

I also had to update my slurm file to get it to go. Unless I had it going with GCC and have it update, I would run into errors with libstdc++, and thus the whole thing would not load at all. Between this is and the error with my nPEx and nPEy, the files really fought back against me.

```

1 #!/bin/bash
2
3 # -
4 # |
5 # | This is a batch script for running a MPI parallel job on Summit
6 # |
7 # | (o) To submit this job, enter: sbatch --export=CODE='/home/trmc7708/Lab2/
    fd_mpi/src' ex01.bat
8 # |
9 # | (o) To check the status of this job, enter: squeue -u <username>
10 # |
11 # -
12

```

```

13 # -
14 # |
15 # | Part 1: Directives
16 # |
17 # -
18
19 #SBATCH --nodes=1
20 #SBATCH --ntasks=4
21 #SBATCH --time=00:01:00
22 #SBATCH --partition=atesting
23 #SBATCH --output=ex01-%j.out
24
25 # -
26 # |
27 # | Part 2: Loading software
28 # |
29 # -
30
31 module purge
32 module update gcc
33 module load gcc
34 module load impi
35
36 # -
37 # |
38 # | Part 3: User scripting
39 # |
40 # -
41
42 echo "=="
43 echo "||"
44 echo "|| Begin Execution of fd in slurm batch script."
45 echo "||"
46 echo "=="
47
48 export OMPI_MCA_opal_common_ucx_opal_mem_hooks=1
49
50 srun -n 4 ./fd -nPEx 2 -nPEy 2 -nCellx 5 -nCelly 5 > tty.out
51
52 echo "=="
53 echo "||"
54 echo "|| Execution of fd in slurm batch script complete."
55 echo "||"
56 echo "=="

```

Listing 2: The Slurm File

## 4 Appendix B: Output

### 4.1 Output of the Slurm

```

1 myPE: 1 x0 = 1
2 myPE: 1 x1 = 2
3 myPE: 1 y0 = 0
4 myPE: 1 y1 = 1

```

```

5 myPE: 2 x0 = 0
6 myPE: 2 x1 = 1
7 myPE: 2 y0 = 1
8 myPE: 2 y1 = 2
9 myPE: 3 x0 = 1
10 myPE: 3 x1 = 2
11 myPE: 3 y0 = 1
12 myPE: 3 y1 = 2
13
14 -----
15
16 F I N I T E   D I F F E R E N C E
17 D E M O   C O D E
18
19 Running on 4 processors
20
21 -----
22
23
24 Input Summary:
25 -----
26 No. PE      in  x-direction: 2
27                y-direction: 2
28 No. Cells in  x-direction: 5
29                y-direction: 5
30
31 myPE: 0 x0 = 0
32 myPE: 0 x1 = 1
33 myPE: 0 y0 = 0
34 myPE: 0 y1 = 1
35
36 Form Linear System:
37 -----
38
39
40 Solve Linear System:
41 -----
42
43 Jacobi/GS converged in 57 iterations
44
45 Plot Results:
46 -----

```

Listing 3: tty.out

## 4.2 Outputs from Alpine

Before switching modules to update gcc and openmpi rather than intel

```

1 ==
2 ||
3 || Begin Execution of fd in slurm batch script.
4 ||
5 ==
6 /home/trmc7708/Lab2/fd_mpi/src/./fd: /usr/lib64/libstdc++.so.6: version 'GLIBCXX_3
  .4.29' not found (required by /home/trmc7708/Lab2/fd_mpi/src/./fd)

```



```

7 /home/trmc7708/Lab2/fd_mpi/src/./fd: /usr/lib64/libstdc++.so.6: version 'GLIBCXX_3
  .4.26' not found (required by /home/trmc7708/Lab2/fd_mpi/src/./fd)
8 srun: error: c3cpu-a5-u34-3: task 0: Exited with exit code 1
9 ==
10 ||
11 || Execution of fd in slurm batch script complete.
12 ||
13 ==

```

Listing 4: Output showing un-updated GCC warnings

After switching:

```

1 Lmod has detected the following error: These module(s) or extension(s) exist
2 but cannot be loaded as requested: "impi"
3   Try: "module spider impi" to see how to load the module(s).
4
5
6
7 ==
8 ||
9 || Begin Execution of fd in slurm batch script.
10 ||
11 ==
12 -----
13 WARNING: There was an error initializing an OpenFabrics device.
14
15   Local host:    c3cpu-a5-u34-3
16   Local device:  mlx5_0
17 -----
18 -----
19 WARNING: There was an error initializing an OpenFabrics device.
20
21   Local host:    c3cpu-a5-u34-3
22   Local device:  mlx5_0
23 -----
24 -----
25 WARNING: There was an error initializing an OpenFabrics device.
26
27   Local host:    c3cpu-a5-u34-3
28   Local device:  mlx5_0
29 -----
30 -----
31 WARNING: There was an error initializing an OpenFabrics device.
32
33   Local host:    c3cpu-a5-u34-3
34   Local device:  mlx5_0
35 -----
36 ==
37 ||
38 || Execution of fd in slurm batch script complete.
39 ||
40 ==

```

Listing 5: Output showing Warnings

### 4.3 Plot Images, showing interconnectivity

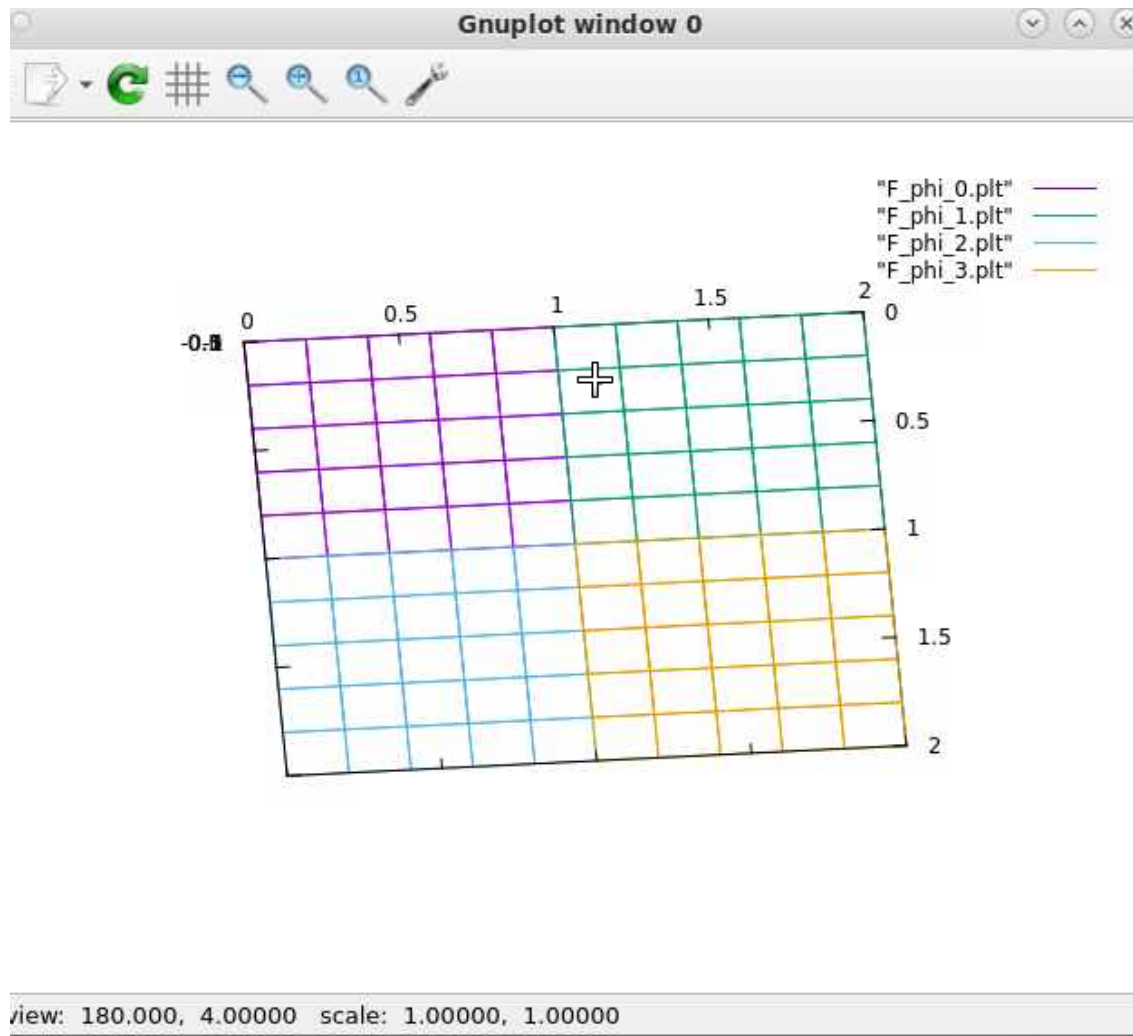


Figure 1: Top Down Image

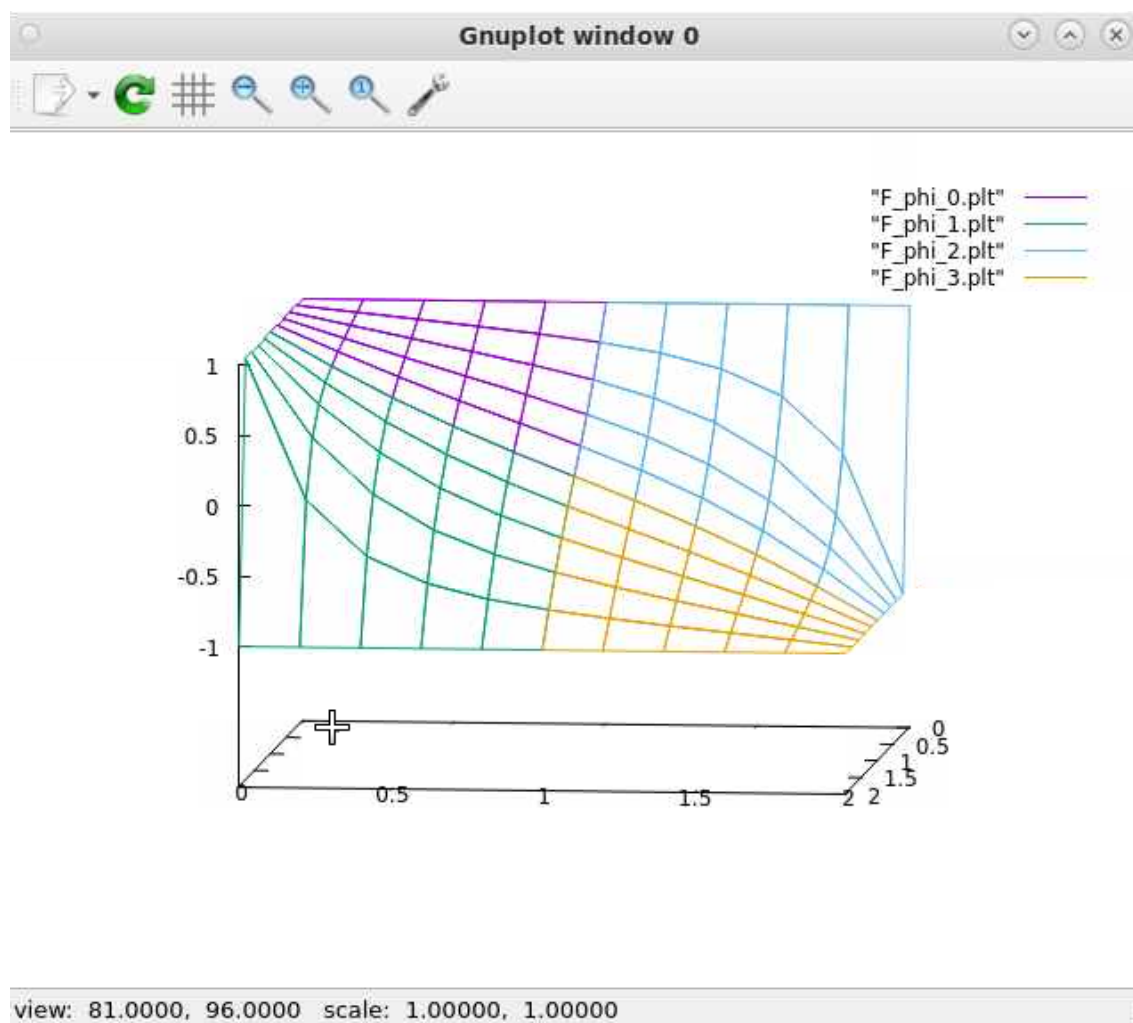


Figure 2: Side Image

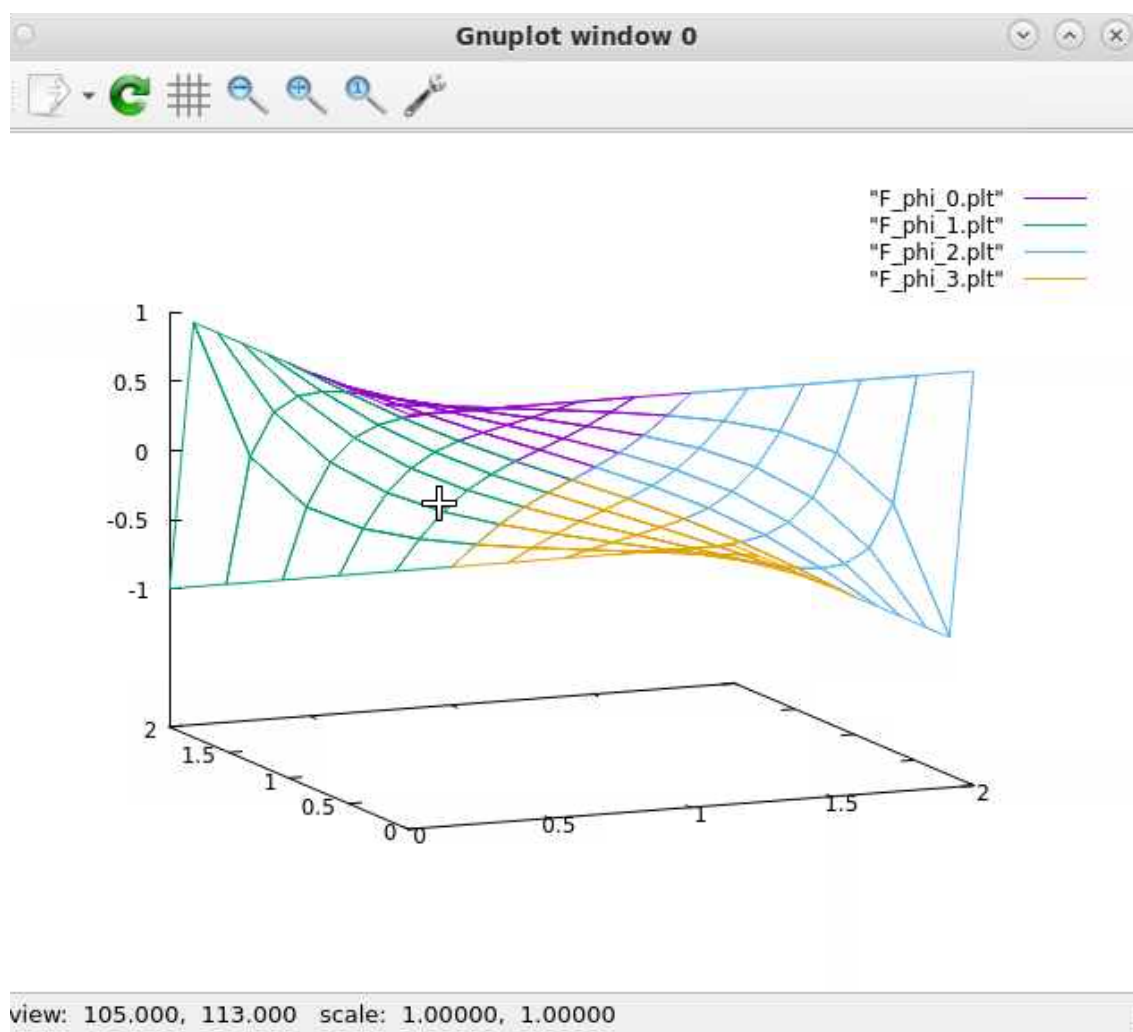


Figure 3: This one just looked cool

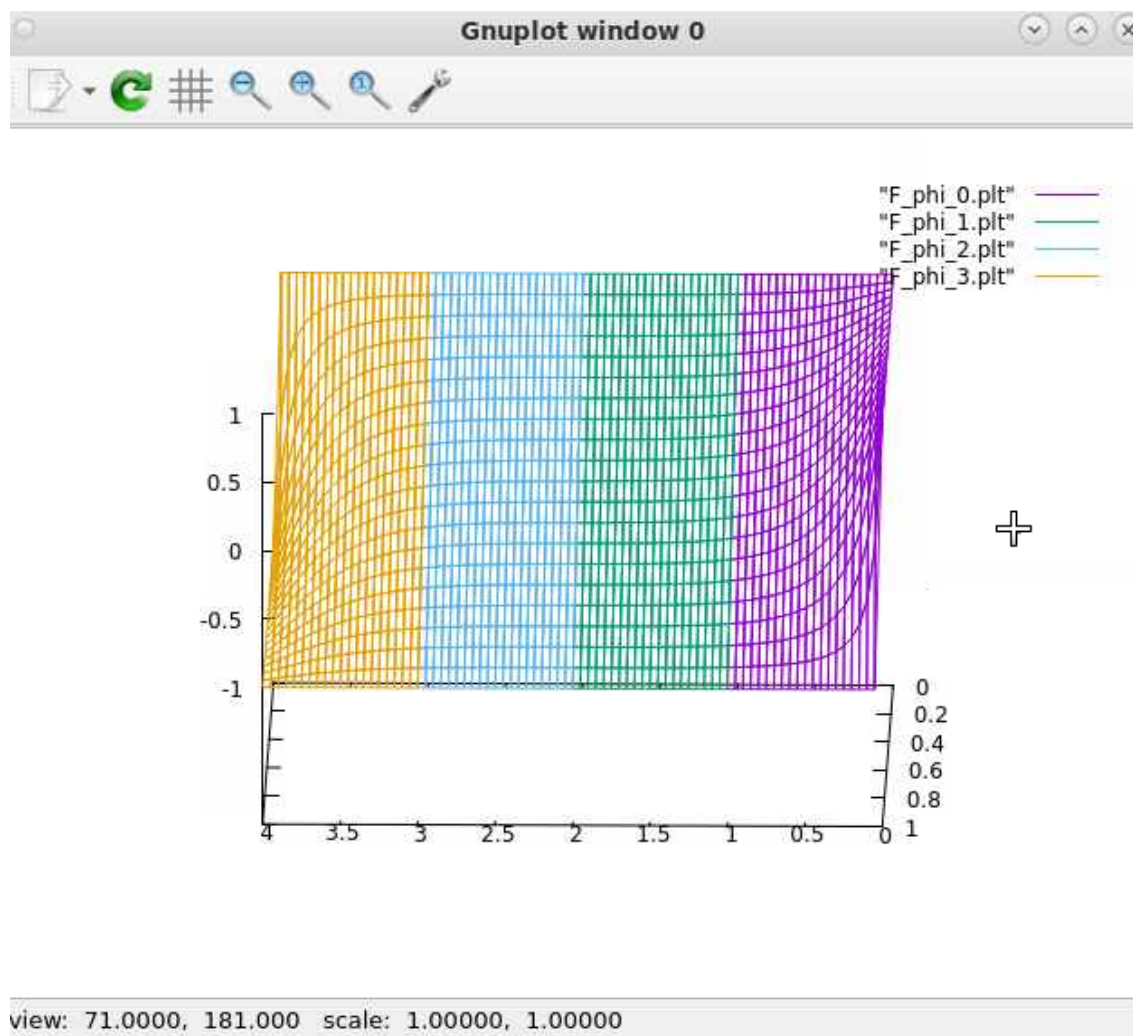


Figure 4: Challenge: nPEx 4 nPEy 1 with 20 cells each