

# HPSC Lab 10

Luna McBride, Noki Cheng

November 17 2023

## 1 Discussion of Approach

The purpose of this lab is to utilize MPIIO to print the mesh into a single file rather than the many files shown off in other labs. This can be done by utilizing the variable phi, which represents the results of the mesh calculations for each specified PE. Phi, however, is a single dimensional vector that is spread out between each of the PEs. PEs<sub>um</sub> has already been run to calculate values on shared nodes luckily enough, so the approach comes down to connecting the different phis on the different PEs and sending them to the shared output file.

This process starts with gathering the phis into a combined phi by using MPI\_Allgather and clipping off the buffer term they all have at the beginning. Phi will be passed to the file as a single dimensional array, so phi must be collected onto a single PE and restructured in a way that the final array outputted resembles the shape of the overall mesh when plotted. The sb python code already handles writing the array into the two-dimensional mesh structure, so this code just needs to weave the phi array so that it can resemble this structure after being sent to sb. This is accomplished using an array called place, while takes note of where each PE's phi starts in the combined phi.

Several for loops are then used to create multiple lines sized by the x axis, with each line connecting as if they were still the contiguous x-PEs shown in a normal mesh. These lines continue until the y-PE boundaries are hit, which then updates place to look at the next set of PEs in the y direction. This results in an array that looks like the mesh, just flipped around. The values of phi in this case go from zero to one from left to right, so the collective phi array needs to be rotated to fit the zero on the bottom and one on the top structure that has existed in the plots from previous labs. This can be done by reversing then transposing the array, utilizing the transpose method discussed here: <https://stackoverflow.com/questions/16737298/what-is-the-fastest-way-to-transpose-a-matrix-in-c> . The transpose function alone resulted with the zeros on top rather than on bottom, so reversing the array first created the expected ones on the top result. The implementation of this methodology was then tested on all possible meshes made by 4 cores (1x1, 1x2, 1x3, 1x4, 4x1, 3x1, 2x1, 2x2) in order to ensure that the algorithm scaled past two PEs in each direction, thus ensuring scalability. The combined phi was then dumped using MPIIO into the required output file, as shown in the below figure.



MPI\_Allgather



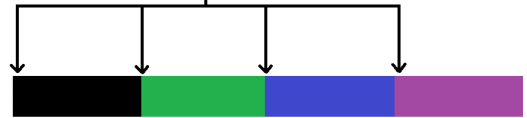
Remove Leading 0's



combined\_phi

# Phi Transformation for MPIIO

place: stores where each phi is located in combined\_phi



Restructure



Structure if it were 2D:



Output



## 2 Results

In this section, we will show the n-to-1 numerical outputs of binary files produced through our use of MPIIO. To avoid redundancy and confusion, some of the configs are fixed

- Solve: jacobi;
- tEnd: 0.15;

- dt: 0.05
- tPlot: 0.05.

We will show three different test cases with different PE and cell sizes, including a special test case 0.

It is also important to note before this, however, that our code uses `MPI_Type_contiguous` instead of `MPI_Type_create_subarray`. This does not allow for the Fortran array style to be used, so the sb code needs to have `nPEy` passed in instead of `nPEx` due to the array being rotated compared to what the program expects.

## 2.1 Test Case 0: Special Test Case

```

1  -----
2
3
4  S O L V E R S
5  D E M O   C O D E
6
7  Running on 4 processors
8
9  -----
10
11
12 Input Summary:
13 -----
14 No. PE      in  x-direction: 2
15                y-direction: 2
16 No. Cells in  x-direction: 3
17                y-direction: 3
18 Linear solver      : jacobi
19 End Time           : 0.15
20 Time Step          : 0.05
21 Plot Interval      : 0.05
22 This is a restart (1/0) : 0
23
24 Plotting at time = 0 Plot ID = 0
25 ***** Writing dump file *****
26 ***** Writing dump file *****
27 ***** Writing dump file *****
28 ***** Writing dump file *****
29 Plotting at time = 0.05 Plot ID = 1
30 ***** Writing dump file *****
31 ***** Writing dump file *****
32 ***** Writing dump file *****
33 ***** Writing dump file *****
34 Plotting at time = 0.1 Plot ID = 2
35 ***** Writing dump file *****
36 ***** Writing dump file *****
37 ***** Writing dump file *****
38 ***** Writing dump file *****
39 Execution Completed Successfully
40

```

```

41 -----
42     sb.py
43 -----
44
45
46 Read user specified file: phi_dump.bin
47 n points x-dir._total    :    7
48
49 Matrix Read from Binary File:
50 -----
51
52    0.0 10.0 20.0   0.1 10.1 20.1 30.1
53    1.0 11.0 21.0   1.1 11.1 21.1 31.1
54    2.0 12.0 22.0   2.1 12.1 22.1 32.1
55    0.2 10.2 20.2   0.3 10.3 20.3 30.3
56    1.2 11.2 21.2   1.3 11.3 21.3 31.3
57    2.2 12.2 22.2   2.3 12.3 22.3 32.3
58    3.2 13.2 23.2   3.3 13.3 23.3 33.3

```

## 2.2 Test Case 1: nPE<sub>x</sub>=4, nPE<sub>y</sub>=1, nCell<sub>x</sub>=nCell<sub>y</sub>=5

```

1  [trmc7708@c3cpu-a2-u15-3 src]$ ./sb_py3.py -f phi_dump.bin -c 5 -n 1
2
3  -----
4     sb.py
5  -----
6
7
8 Read user specified file: phi_dump.bin
9 n points x-dir._total    :    6
10
11 Matrix Read from Binary File:
12 -----
13
14 1.000 1.000 1.000 1.000 1.000 1.000
15 0.950 0.946 0.943 0.943 0.946 0.950
16 0.900 0.892 0.887 0.887 0.892 0.900
17 0.850 0.838 0.831 0.831 0.838 0.850
18 0.800 0.785 0.775 0.775 0.785 0.800
19 0.750 0.732 0.720 0.720 0.732 0.750
20 0.700 0.679 0.666 0.666 0.679 0.700
21 0.650 0.627 0.613 0.613 0.627 0.650
22 0.600 0.575 0.560 0.560 0.575 0.600
23 0.550 0.525 0.509 0.509 0.525 0.550
24 0.500 0.474 0.458 0.458 0.474 0.500
25 0.450 0.425 0.409 0.409 0.425 0.450
26 0.400 0.376 0.361 0.361 0.376 0.400
27 0.350 0.327 0.313 0.313 0.327 0.350
28 0.300 0.279 0.267 0.267 0.279 0.300
29 0.250 0.232 0.221 0.221 0.232 0.250
30 0.200 0.185 0.176 0.176 0.185 0.200
31 0.150 0.138 0.131 0.131 0.138 0.150
32 0.100 0.092 0.087 0.087 0.092 0.100
33 0.050 0.046 0.044 0.044 0.046 0.050
34 0.000 0.000 0.000 0.000 0.000 0.000

```

## 2.3 Test Case 2: nPE<sub>x</sub>=2, nPE<sub>y</sub>=2, nCell<sub>x</sub>=nCell<sub>y</sub>=3

```
1 [trmc7708@c3cpu-a2-u15-3 src]$ ./sb_py3.py -f phi_dump.bin -c 3 -n 2
2
3 -----
4     sb.py
5 -----
6
7
8 Read user specified file: phi_dump.bin
9 n points x-dir._total    :    7
10
11 Matrix Read from Binary File:
12 -----
13
14 1.000 1.000 1.000 1.000 1.000 1.000 1.000
15 0.833 0.822 0.814 0.811 0.814 0.822 0.833
16 0.667 0.647 0.633 0.628 0.633 0.647 0.667
17 0.500 0.478 0.462 0.456 0.462 0.478 0.500
18 0.333 0.314 0.301 0.296 0.301 0.314 0.333
19 0.167 0.156 0.148 0.145 0.148 0.156 0.167
20 0.000 0.000 0.000 0.000 0.000 0.000 0.000
```

## 3 Self Evaluation

### 3.1 Luna

This assignment was split in two by who updated phi and who got the output to run. I updated phi, and may have gone a little far with the concept. The updates became a game of questions like "but what if the mesh is expanded?" and "what if the transpose utility does not work on a mesh that is not square?". Allgather also only worked correctly if phi was an array and not a vector, so add a conversion array to go with the flipping array needed for transpose and the array to hold the data once the phis were weaved together. There was very likely another way to update phi that reduced space and time complexity, but I did not want to get in the same loop of "try something that would work great in Python with NumPy/Pandas but end up empty handed in C++" like happened in the previous lab. The code works and accounts for meshes larger than 2x2, so this is plenty.

### 3.2 Noki

The lab overall went well, and we successfully implemented the n-to-1 parallel output using MPIIO. The process of adapting the MPIIO code to handle a 1-D array was challenging but educational. Understanding and addressing the overlapping nodes at PE boundaries required careful consideration, which is mostly contributed by Luna.

- **Hard Parts:**

- Adapting MPIIO for 1-D output: Modifying the MPIIO code to work with a 1-D array instead of a 2-D matrix posed a significant challenge. Ensuring correct indexing and avoiding duplication of PE boundary values required careful debugging.

- Addressing PE Boundary Overlaps: The transientDiffusion\_1mat code had overlapping nodes at PE boundaries, necessitating thoughtful modifications to the MPIIO approach. This required a deep understanding of both the finite difference code and MPIIO.

- **What Went Well:**

- Successful Parallel Output: The implemented solution demonstrated the desired n-to-1 parallel output capability on a 2x2 and 4x1 PE decomposition without any evident artifacts of the number of processes.
- Clear Test Case Verification: The test case, showcasing the i-j location and PE residence of phi values, verified the correctness of the n-to-1 writer. The results from sb\_py3.py aligned with expectations.

In conclusion, while the lab presented challenges, the successful implementation of the parallel output functionality and the clear verification through test cases indicate a satisfactory completion of the task.

## 4 Appendix A: C++ Source Code

```

1 //
  =====
2 //  ||
  ||
3 //  ||      transientDiffusion
  ||
4 //  ||      -----
  ||
5 //  ||      T R A N S I E N T D I F F U S I O N
  ||
6 //  ||
  ||
7 //  ||      D E M O N S T R A T I O N   C O D E
  ||
8 //  ||      -----
  ||
9 //  ||
  ||
10 //  ||      Developed by: Scott R. Runnels, Ph.D.
  ||
11 //  ||      University of Colorado Boulder
  ||
12 //  ||
  ||
13 //  ||      For: CU Boulder CSCI 4576/5576 and associated labs
  ||
14 //  ||
  ||
15 //  ||      Copyright 2020 Scott Runnels
  ||

```

```

16 //  ||
17 //  ||                               Not for distribution or use outside of the
18 //  ||                               this course.
19 //  ||
20 //  ||
21 //  =====
22 #include <mpi.h>
23 #include <cstdlib>
24
25 //  =====
26 //  ||
27 //  ||  M P I I O   D u m p   W r i t e r
28 //  ||
29 //  ||
30 //  =====
31 void write_mpiio_dump(mpiInfo &myMPI)
32 {
33
34     printf("***** Writing dump file *****\n");
35
36     int nRealx  = myMPI.nRealx;    // Number of real points in the x- and y-
37     int nRealy  = myMPI.nRealy;    // directions
38     int myPE, numPE;    // This process' ID and the number of processes total
39     int nPEx, nPEy;    // Number of processes in the x- and y-directions, in the 2D
40     int iPE, jPE;    // grid of processes
41
42     // This process' logical ID in the 2D array of processors
43     nPEx = myMPI.nPEx;
44     nPEy = myMPI.nPEy;
45     myPE = myMPI.myPE;
46     numPE = myMPI.numPE;
47     jPE = int(myPE/nPEx);
48     iPE = myPE - jPE*nPEx;
49
50     // float *phiVal;
51     // phiVal = new float [nField+1];
52     // for (int i = 1; i <= nField; i++)
53     //     phiVal[i] = i + myPE/10.;
54
55     MPI_Barrier(MPI_COMM_WORLD);
56     int num_send = phi.size();

```

```

56  int max_send;
57
58  MPI_Allreduce(&num_send, &max_send, 1, MPI_INT, MPI_MAX, MPI_COMM_WORLD);
59
60  int gather_size = max_send*myMPI.numPE;
61  double* phi_array = new double[phi.size()];
62  int* pe_gather = new int[phi.size()];
63
64  for(int i = 0; i < phi.size(); ++i){
65      phi_array[i] = phi[i];
66      pe_gather[i] = myMPI.myPE;
67  }
68
69  double* phiTest = new double[gather_size];
70  int* phiPE = new int[gather_size];
71
72      for ( int i = 0 ; i < gather_size ; ++i ) {phiTest[i] = 0.; phiPE[i] = -1;}
73
74  MPI_Allgather( pe_gather , max_send , MPI_INT , phiPE , max_send , MPI_INT ,
75                MPI_COMM_WORLD);
76  MPI_Allgather( phi_array , max_send , MPI_DOUBLE, phiTest , max_send ,
77                MPI_DOUBLE, MPI_COMM_WORLD);
78
79  MPI_Barrier(MPI_COMM_WORLD);
80  //if(myMPI.myPE==0){
81  double* combined_phi = new double[gather_size - myMPI.numPE];
82  int prev_pe = -1;
83  int phiCount = 0;
84
85  //Removes the padding 0 from the front and puts all actual values into a
86  combined phi variable
87  for(int i = 0; i < gather_size; ++i){
88      if(prev_pe != phiPE[i]) prev_pe = phiPE[i];
89      else {
90          combined_phi[phiCount] = phiTest[i];
91          phiCount++;
92      }
93  }
94
95  int* place = new int[myMPI.numPE];
96  for (int i = 0; i < myMPI.numPE; ++i){
97      place[i] = 0 + (phi.size()-1)*i;
98  }
99
100  int big_mesh_x = myMPI.nRealx + nCellx*(myMPI.nPEx-1);
101  int big_mesh_y = myMPI.nRealy + nCelly*(myMPI.nPEy-1);
102
103  int current_pe = 0;
104  int pe_change = 0;
105  int current_position = 0;
106  int num_changedy = 0;
107  int num_changedx = 0;
108
109  double* fixed_phi = new double[big_mesh_x*big_mesh_y];
110  double* flipped_phi = new double[big_mesh_x*big_mesh_y];

```



```

108 int i = 0;
109
110 for (int y = 0; y < big_mesh_y; ++y){
111     if(y%nCelly==0 && y!=0 && (num_changedy+1 < myMPI.nPEy)){
112         pe_change += myMPI.nPEx;
113         num_changedy++;
114     }
115     current_pe = current_pe + pe_change;
116
117     for(int x = 0; x < big_mesh_x; x++){
118         if(x%nCellx==0 && x!=0 && (num_changedx+1 < myMPI.nPEx)){
119             place[current_pe]++;
120             current_pe++;
121             num_changedx++;
122         }
123
124         current_position = place[current_pe];
125         fixed_phi[i] = combined_phi[current_position];
126         i++;
127         place[current_pe]++;
128     }
129     current_pe = 0;
130     num_changedx = 0;
131 }
132 int c = 0;
133 //Reverse the array so that it can transpose to the correct form
134 double temp = 0.;
135 int mesh_len = (big_mesh_x*big_mesh_y);
136 for(int n = 0; n < mesh_len/2; n++) {
137     temp = fixed_phi[mesh_len-n-1];
138     fixed_phi[mesh_len-n-1] = fixed_phi[n];
139     fixed_phi[n] = temp;
140 }
141
142 //Transpose the phi so that it can look correct
143 //Source for transpose: https://stackoverflow.com/questions/16737298/what-is-the-fastest-way-to-transpose-a-matrix-in-c
144 for(int n = 0; n < big_mesh_x*big_mesh_y; n++) {
145     int i = n/big_mesh_y;
146     int j = n%big_mesh_y;
147     flipped_phi[n] = fixed_phi[big_mesh_x*j + i];
148 }
149
150 //print out phi
151 <<<<<< HEAD
152 //if (myMPI.myPE==0){
153     //cout << "[" << endl;
154     //for(int n = 0; n < big_mesh_x*big_mesh_y; n++) {
155         //if(n%big_mesh_y==0) cout << endl;
156         //cout << flipped_phi[n] << " ";
157         //}
158     ///cout << "]" << endl;
159     //
160
161 //if(myMPI.myPE == 0){

```

```

162 MPI_Datatype myRealNodes;
163
164 MPI_Type_contiguous(nTotal, MPI_DOUBLE, &myRealNodes);
165
166 MPI_Type_commit(&myRealNodes);
167 //
168 //
    -----

169
170 MPI_File_delete( "phi_dump.bin", MPI_INFO_NULL );
171
172 MPI_File fh;
173
174 MPI_File_open(MPI_COMM_WORLD, "phi_dump.bin", MPI_MODE_CREATE | MPI_MODE_WRONLY,
    MPI_INFO_NULL, &fh);
175
176 MPI_File_set_view (fh, 0, MPI_DOUBLE, myRealNodes, "native", MPI_INFO_NULL);
177 //                                     |                                     |
178 //                                     |                                     |
179 //                                     The data type of each element         Could contain
    optimization hints,                                                         e.g., striping
180 //                                     |                                     |
181 //                                     |                                     |
    -----

182
183 // (8) Perform the collective write operation
184
185
186 //
    -----

187 //                                     The sub_array describing this
188 // File pointer                         PE's real-nodes part of A
189 //                                     |
190 //                                     |
191 if(myMPI.myPE==0){
192 MPI_File_write(fh, flipped_phi, 1, myRealNodes, MPI_STATUS_IGNORE);
193 }
194 //                                     |
195 //                                     |
196 // Memory location of original
197 // data set for which the
198 // sub arrays were created.
199 //
    -----

200
201 MPI_File_close(&fh);
202 MPI_Type_free(&myRealNodes);
203 //}
204 free(pe_gather);
205 free(phiPE);

```

```

206 free(phi_array);
207 free(phiTest);
208 free(combined_phi);
209 free(fixed_phi);
210 free(flipped_phi);
211
212 //}
213 MPI_Barrier(MPI_COMM_WORLD);
214 }

```

## 5 Appendix B: sb\_py3.py Script

```

1  #!/usr/bin/python3
2
3  import os
4
5  import sys
6  import getopt
7  import glob
8  import struct
9
10
11 # ==
12 # ||
13 # || Main Program
14 # ||
15 # ==
16
17 def sb(argv):
18
19     os.system('clear')
20
21     print()
22     print( '-----')
23     print( '    sb.py')
24     print( '-----')
25     print()
26
27     # -
28     # |
29     # | Get user arguments
30     # |
31     # -
32
33     readFile = 'NULL'
34     nCellx   = 0
35     nPEx     = 0
36
37     try:
38         opts, args = getopt.getopt(argv,"h f: c: n:")
39
40     except:
41         print( 'Error in command-line arguments.')
42         exit(0)
43

```

```

44     for opt, arg in opts:
45
46         if opt == "-h":
47             print( './sb_py3.py -f <binary file to be read> -c <no. cells in x/y
dir *per PE*> -n <nPEx>')
48             exit(0)
49
50         if opt == "-c":
51             nCellx = int(arg)
52
53         if opt == "-n":
54             nPEx = int(arg)
55
56         elif opt == "-f":
57             readFile = arg
58
59     # -
60     # |
61     # | Check user arguments
62     # |
63     # -
64
65     if readFile == 'NULL':
66         print("Fatal Error: Specify binary file with -f")
67         exit(0)
68
69     if nCellx == 0:
70         print("Fatal Error: Specify no. cells in x/y direction with -c")
71         exit(0)
72
73     if nPEx == 0:
74         print("Fatal Error: Specify no. PE in x direction with -n")
75         exit(0)
76
77     # -
78     # |
79     # | Specifications for reading the binary and writing to tty
80     # |
81     # -
82
83     nx_total      = nPEx*nCellx + 1
84     DecimalSize   = 8          # 4 for single precision (float) 8 for double
precisioni (double)
85     DecimalFormat = 'd'       # 'f' for single, 'd' for double
86
87     # -
88     # |
89     # | Read binary file and print it to screen
90     # |
91     # -
92
93     print()
94     print( 'Read user specified file: ' + readFile)
95     print( 'n points x-dir._total   : ' , nx_total)
96

```

```

97     file = open(readFile, "rb")
98
99     print()
100    print( "Matrix Read from Binary File:")
101    print( "-----")
102    print()
103
104    # Read first data
105
106    Num = file.read(DecimalSize)
107
108    count_x_total = 0
109
110    # Continue reading data until end of file
111
112    while Num:
113
114        # Read and print data with no line feed
115
116        floatValue = struct.unpack(DecimalFormat,Num)[0]    # 'f' for single, 'd'
for double
117        print( "{:_.3f}".format(floatValue),end=' ')
118
119        # Put line feed at the end of the mesh
120
121        count_x_total += 1
122        if count_x_total == nx_total:
123            print()
124            count_x_total = 0
125
126        Num = file.read(DecimalSize)
127
128
129    file.close()
130
131
132
133
134    print()
135    print()
136
137
138
139 if __name__ == "__main__":
140     sb(sys.argv[1:])

```